

S-043/Stat-151

Analysis for Clustered and Longitudinal Data
(Multilevel & Longitudinal Models)


Multilevel Logistic Regression

R&B Chapter 10

Also Chapter 14 (mainly 14.2 and 14.3) of G&H

Unit 5: Where are we?

We are here



Today's Goals

Quickly review some highlights of Poisson regression

Move on to Logistic regression... but with multiple levels.

- ★ We will look at a new dataset about kids being retained a year
- ★ We will see how we can describe variation in retention rates across schools
- ★ We will also see how we can model individual differences (e.g., between girls and boys)

By the end you will be able to

- ★ Fit a multilevel generalized model
- ★ Understand that the estimated coefficients need to be transformed to make them interpretable
- ★ Use predict() to make the results of a logistic model more interpretable.

Back to the Poisson GLM

Super Cinder (aka Cinder) wants us to (1) look awesome for Halloween, and (2) revisit last lecture for a moment...



“I love
wearing a
costume.”

The anatomy of GLMs via Poisson

Our Covariates

$$X_i = (1, Eth_i, Pcnt_i)$$



Each observation is a bunch of covariates plus an observed outcome (a count) plus an exposure.

Our Linear Predictor

$$\eta_i = X'_i \beta$$

“eta”



Our rate is a transform of our linear predictor


$$\theta_i = g(\eta_i) = \exp(X'_i \beta)$$

exposure X rate gives how many events we *expect* to see



$$\hat{Y}_i = u_i \theta_i = \exp(\log(u_i) + X'_i \beta)$$

We can shove our exposure into the linear predictor, making it look like a covariate.



This is our inverse “link function”

We **actually** observe a Poisson draw around what we expected

$$Y_i = \text{Poisson}(u_i \theta_i)$$

(Our residual is then the difference)

$$resid_i = Y_i - \hat{Y}_i$$



$$E[Y_i] = \hat{Y}_i, sd(Y_i) = \sqrt{\hat{Y}_i}$$

This step is where the randomness is!



Our fit model revisited

```
> fit.4 <- glm (stops ~ factor(eth) + factor(precinct),  
family=quasipoisson, Just think of quasipoisson as  
Poisson, fixed.  
          offset=log(past.arrests), data=stops )  
> display(fit.4)
```

	coef.est	coef.se
(Intercept)	-1.38	0.24
factor(eth) 2	0.01	0.03
factor(eth) 3	-0.42	0.04
factor(precinct) 2	-0.15	0.35
factor(precinct) 3	0.56	0.27
...		
factor(precinct) 74	1.15	0.27
factor(precinct) 75	1.57	0.35

n = 225, k = 77
residual deviance = 3427.1, null deviance = 46120.3
(difference = 42693.1)
overdispersion parameter = 21.9

We can believe these standard errors now.

(And they are substantially larger.)





Pop! Quiz

$$\text{logit} = \text{log odds} = \log\left(\frac{p}{1-p}\right) = \alpha$$

$$\text{In reverse: } \text{logit}^{-1}(\alpha) = \frac{1}{1 + \exp(-\alpha)} = \frac{\exp(\alpha)}{1 + \exp(\alpha)}$$

If the log odds is 3, what is the chance of winning?



Multilevel &
Logistic,
here we come!

Today's Data: Thailand Retention Data

Students in schools, as per usual.

Students are either retained or not (*binary outcome*)

Variables:

From R&B Chapter 10

★ Retained (1067 of 7516 retained) This is our outcome.

★ Gender

★ Pre-primary school experience. (Whether they have prior experience with school.)

★ Mean SES of school


Goal: (1) describe variation in retention rates across schools, and

(2) describe how students with different prior experiences differ.


(3) see if school mean SES is associated with rate of retention.




Distribution of schools and their retention rates



Retention probabilities for four groups



Higher Mean SES is associated with lower retention



Components of a Generalized Linear Model

$$Y_{ij} \sim F(\mu_{ij}, \psi)$$

$$E[Y_{ij}] = \mu_{ij}$$

Each observation has some predicted mean (the μ_{ij} , often called \hat{Y}_{ij})

We specify some **sampling distribution** for our actual outcomes Y_{ij} around the μ_{ij}

F is a distribution (e.g., Normal, Poisson, Binomial, etc., that takes the mean and possibly other parameters such as a variance (the ψ).

Component: Level-1 Link Function

"Eta"  $\eta_{ij} = g(\mu_{ij})$ with $Y_{ij} \sim F(\mu_{ij}, \psi)$

Eta is our **linear predictor**

We tie our linear predictor to the *actual* prediction
with our **link function**

So far, we don't have any multilevel anything going on...

We have a linear model to predict our eta for each unit. Then we transform eta to get our actual prediction mu.

Component: Level 1 Structural Model

Note how eta is our linear predictor—
and it looks just like classic OLS

$$\eta_{ij} = \beta_{0j} + \beta_{1j}X_{1ij} + \beta_{2j}X_{2ij}$$

Note lack of residual term. (Where did it go?)

Now we see grouping structure: Each coefficient could vary by group, as shown by the j subscripts.

To denote pooling across all groups, omit the j .

$$\eta_{ij} = \beta_0 + \beta_1 X_{1ij} + \beta_2 X_{2ij}$$

Q: What kind of multilevel model is this?

Component: The Level 2 Model

Each level 1 coefficient that varies by group has a level-2 equation:

$$\beta_{qj} = \gamma_{q0} + \gamma_{q1} W_{1j} + \dots + u_{qj}$$

Our level-2 is (usually) in our classic OLS framework (i.e., this is exactly the same as we have seen before).

Each coefficient in our level-1 model is modeled as

- ★ a linear combination of level-2 predictors with level-2 coefficients.
- ★ a random level-2 noise term (the random effect)

Logistic Binomial Model

If we have m_{ij} observations for unit i in group j , and Y_{ij} is the number of successes, we have a random intercept model of

$$Y_{ij} \sim \text{Binom}(m_{ij}, \mu_{ij})$$

m_{ij} coin flips, each with μ_{ij} chance of heads

$$E[Y_{ij}] = m_{ij}\mu_{ij}$$

**We have $m_{ij} * \mu_{ij}$ because m observations each with μ_{ij} chance
(The m_{ij} is like the exposure from the Poisson model.)**

$$\mu_{ij} = g^{-1}(\eta_{ij}) = \text{logit}^{-1}(\eta_{ij}) = \frac{1}{1 + \exp(-\eta_{ij})}$$

logistic link function

$$\eta_{ij} = \beta_{0j} + \beta_1 X_{ij}$$

linear predictor has an intercept varying by group

$$\begin{aligned} \beta_{0j} &= \gamma_{00} + u_{0j} \\ u_{0j} &\sim N(0, \tau) \end{aligned}$$

**Level 2 is same-old, same-old. Only level 1 has the links
and fancy connections to outcomes**

This is a random intercept model with a logistic model for level 1



The Simplest Model

```
> M0 = glmer( REPEAT ~ 1 + (1|SCHOOLID), data=dat,  
+               family=binomial(link="logit") )
```

Warning messages:

```
1: In checkConv(attr(opt, "derivs"), opt$par, ctrl = control  
$checkConv, :  
  Model failed to converge with max|grad| = 0.184618 (tol = 0.001,  
component 1)  
2: In checkConv(attr(opt, "derivs"), opt$par, ctrl = control  
$checkConv, :  
  Model is nearly unidentifiable: very large eigenvalue  
  - Rescale variables?
```

```
> display( M0 )  
glmer(formula = REPEAT ~ 1 + (1 | SCHOOLID), data = dat, family =  
binomial(link = "logit"))  
coef.est  coef.se  
-2.23     0.00
```

Error terms:

Groups	Name	Std.Dev.
SCHOOLID	(Intercept)	1.29
Residual		1.00

```
number of obs: 7516, groups: SCHOOLID, 356  
AIC = 5547, DIC = 4087  
deviance = 4815.1
```

! We are not converging. We are also getting different answers from R&B. Not sure what is going on.

What does the intercept mean?

A “Typical” school has a random intercept of 0, meaning we have:

- ★ Expected log-odds of retention = -2.23
- ★ I.e., odds of $\exp(-2.23)$ = 0.109 or about 1 to 9
- ★ I.e., probability of $1/(1+\exp(2.22)) = 0.097$

Compare with overall rate of retention of 14%


KEY POINT:

“Typical” is not the same as population average due to non-linearity!

Assessing cross-school variation (Also, ICC? Not so useful.)

Variation is **immense**.

Some schools have retention rates of near 0% and some are above 50%!



What about ICC?

- ★ ICC is ratio of level-2 variance to total variance.
- ★ But different schools have different kids, and so different variances due to Binomial.
- ★ There is no natural ratio here.
- ★ Instead draw pictures and show people **how much schools vary**.



Mean SES (Centered)

```
> M1 = glmer( REPEAT ~ 1 + SEX + PPED + MSESC + (1|SCHOOLID),  
+               data=dat,  
+               family=binomial(link="logit") )
```

```
> display( M1 )
```

	coef.est	coef.se
(Intercept)	-2.24	0.11
SEXboy	0.53	0.08
PPEDyes	-0.63	0.10
MSESC	-0.30	0.21

Error terms:

Groups	Name	Std.Dev.
SCHOOLID	(Intercept)	1.27
	Residual	1.00

number of obs: 7516, groups: SCHOOLID, 356
AIC = 5456.9, DIC = 4014.4
deviance = 4730.7



What is the probability of a boy with no prior experience in an average school being retained?



Our covariates:
Gender
pre-school experience
mean school SES

What levels?



Convert the log to odd multipliers **(And examining our impact of Mean SES)**

```
> exp( fixef( M1 ) )
(Intercept) 0.107
```

SEXboy
1.706

PPEDyes
0.534

MSESC
0.743

The odds of a boy being retained is 1.7 times those of a girl, all other things being equal (including school).

For schools, each unit increase in Mean SES corresponds to a reduction of x0.74

```
> sd( sdat$mean.ses)
[1] 0.381
```

so 1 SD above average gives:

```
> exp( 0.381 * 0.743 ) = 0.89
```

i.e., schools 1 SD above average have 90% of the retention rate (on average)



Predicting individual retention

```
> dat$eta.hat = predict( M1 )  
> dat$eta.hat = predict( M1, type = "link" )  
> sample( dat$eta.hat, 7 )  
[1] -0.91311147 -3.98446224 -0.86143874 -0.05038648  
-3.16325990 -2.55520422 -2.68551484
```

Two ways, same thing

```
> dat$mu.hat = predict( M1, type="response" )  
> sample( dat$mu.hat, 7 )  
[1] 0.11416657 0.45092846 0.46885009 0.02925308  
0.15426398 0.20458749 0.10572866
```

This is the linear predictor

These are probabilities!

**We can predict our linear predictors
for our individuals, or predict
our outcome (response).**

This is using the link function or not.

Recap of our research questions


RQ1: describe variation in retention rates across schools, and

RQ2: describe how students with different prior experiences differ.


RQ3: see if school mean SES is associated with rate of retention.

- ★ Look at coefficient
- ★ Make plot to assess impact


Boys and girls across schools



Our four groups of children (for different schools)



Impact of Mean SES



Population Average vs. Individual School

So far we are predicting for individuals in individual schools.
As things are non-linear, this is not the same as estimating averages across the population.

E.g., if we are investigating pre-primary **experience** and repetition we have two possible questions:

- A: What is the association between pre-primary experience and risk of repetition at a given school?
- B: What is the overall association of these two things averaged across schools?

*The coefficient in our model is **NOT** the answer to the second question.*



Fitting a Linear Model instead...

```
> M1.lin = lmer( as.numeric(REPEAT == "yes") ~ 1 + SEX +
PPED + MSESC + (1|SCHOOLID), data=dat )
> display( M1.lin )
lmer(formula = as.numeric(REPEAT == "yes") ~ 1 + SEX + PPED +
+ MSESC + (1 | SCHOOLID), data = dat)
            coef.est    coef.se
(Intercept)  0.15      0.01
SEXboy       0.05      0.01
PPEDyes     -0.06      0.01
MSESC        -0.03      0.02
```

Error terms:

Groups	Name	Std.Dev.
SCHOOLID	(Intercept)	0.14
	Residual	0.32

```
number of obs: 7516, groups: SCHOOLID, 356
AIC = 4724.2, DIC = 4655
deviance = 4683.5
```



...has some trouble

```
# our predictions for all students  
pd = predict( M1.lin )
```

```
skimr::skim( pd )
```

Skim summary statistics

some predictions are below 0, indicating negative chance of retention.


```
— Variable type: numeric  
variable missing complete n mean sd p0 p25 p  
pd 0 7516 7516 0.14 0.13 -0.065 0.045 0.
```

```
> mean( pd <= 0 )  
[1] 0.05361895
```

about 5% of our predictions are negative!

1 Also, the sd of the random intercept is 0.14, indicating entire schools with negative retention

Recap



Pepper the Black Cat

What we have seen today

<http://cs179.org/lec52>

Multilevel data with binary outcomes can be thought of as:

- ★ A logistic regression framework at level 1
- ★ A vanilla random-effects world at level 2 (no logistic stuff here!)

We can do this because level 2 equations are predicting the *coefficients* of the level 1 regression models, and these coefficients are continuous.

Interpreting models with nonlinear links is tricky.
Solution: make plots and describe trends and variation.

Appendix: Red States vs. Blue States Another example of multilevel logistic regression

See Chapter 14.2 of G&H
Interesting case study, worth a
careful read for interpreting
models, etc.

Red States vs. Blue States


The two following things are true:

- ★ Richer states tend to have a higher vote share for Democratic candidates
- ★ Richer people tend to have a higher vote share for Republican candidates


We can model this to understand what is going on with this seeming “paradox.”

(See A Gelman’s “Red State, Blue State, Rich State, Poor State:
Why Americans Vote the Way They Do”³⁴)

Slope vs. state income, 2000



Varying-intercept model, 2000



Varying intercept, varying slope

$$Pr[y_i = 1] = \text{logit}^{-1}(\alpha_j[i] + \beta_j[i]x_i)$$

$$\alpha_j = \gamma_{00} + \gamma_{01}u_j + e_{0j}$$


$$\beta_j = \gamma_{10} + \gamma_{11}u_j + e_{1j}$$

Level 2 is the same old, same old.


We have joint normally distributed residuals at level 2 as per usual.

It is just level 1 that gets the extra structure with links, etc.


Varying-intercept, varying-slope model, 2000



Intercept vs. state income, 2000



Rich people vote republican




Each year vote for republican regressed onto income.

We see the coefficient (slope) of the model for each year.

0 would be no relationship.

Rich states vote democrat



Regression of republican vote share on average state income.

Right two plots broken down by north & south

**Model fit separately for
each election year.**

Take-away messages

Interpreting the model is where the work is

You should rely on visualizations and descriptions of different “typical cases”

(E.g., a rich, middle, and poor state)

You should plot the individual random effects against other level-2 covariates to look for trends

Describe across the groups, but avoid making strong claims about specific groups.